


Flash Storage

1. From reviews:
 - a. SE-Util vs SE-Merge
 - b. Why does write back perform better while write through can use silent eviction?
 - c. Why different data vs log blocks?
2. Flash device basics
 - a. All flash:
 - i. Stores charge persistently that can later be sensed
 - ii. Can only write changing a 1 to a zero; changing a zero to a 1 requires a coarse-grained and slow erase operation
 - b. NOR flash:
 - i. Allows random access (can be accessed like ROM)
 - ii. Fast – 75ns, close to DRAM speeds for read, slow writes, very slow for erase (1 second)
 - iii. Used for storing system code – can execute in place without copying to DRAM first, saving memory
 - iv. Lower density than NAND
 - c. NAND flash
 - i. Page access only as a device; cannot be accessed as memory
 - ii. 25 us reads, 1.5 ms erase, 100us writes
 - d. All have endurance issues
 - i. After erasing too many times, stops working.
3. Flash packaging
 - a. Embedded flash
 - i. No controller; raw flash accessible to CPU
 - ii. All management (garbage collection, striping of data) must be done by OS
 - iii. Used in embedded devices
 - b. Removable flash – CompactFlash, SD cards
 - i. Simple controller in the package
 - ii. Depending on generation:
 1. Simple mapping: writes do a read/modify/write in place
 2. Limited wear leveling
 - iii. Used for consumer devices with sequential workloads
 - c. SSDs – managed flash
 - i. Controller implements translation layer


- ii. Controller + software manages parallelism remapping, garbage collection
- 4. History of storage technologies
 - a. Disks remain and will pretty much always be larger and slower
 - b. Many candidate technologies- e.g. SD cards, MEMS
 - i. All faster than disks, but smaller and more expensive?
 - c. Why flash?
 - i. Found a use disks couldn't work
 - 1. low power important; disks are ~5 watts
 - ii. Capacity not as important
 - 1. Cell phones, early phones & mp3 players
 - 2. How much is a 8 GB hard disk compared to 8GB flash?
 - a. answer: more
 - iii. Market was big enough to sustain development and make money
 - iv. Once Flash was mature, people started building SSDs for PCs/servers
 - d. Nothing else has done this yet but disks
 - e. Problem with adopting new storage technologies:
 - i. Show memory/disk curve – speed vs capacity
 - ii. For large capacity, disk always cheapest
 - iii. For high performance, memory best
 - iv. Hard to live in the middle



Take Away

Market compares on \$/GB, \$/Performance

v.
vi.




Take Away

Displacement possible when better in both dimensions

- vii. © 2003 S. B. H. Lee @ Georgia Tech. All Rights Reserved.
- viii. Dead zone – at top; where more expensive than L1, and on right, where same for L2


The Cache Zone



Take Away

Cache Zone storage increases system cost

- ix.
- 5. SSD Organization
 - a. Physical layout



- i.
- ii. Multiple dimensions: dies, split into planes, with blocks containing pages
- iii. Parallelism across planes, between some dies internally
- iv. SLC vs MLC:
 1. SLC is one bit (2 voltage levels) per cell, fast
 2. MLC is 2 bits (4 voltage levels) per cell: higher capacity, lower speed, lower endurance (100,000 erases per SLC, 10K for MLC, 1 K for TLC)
- v. Connections:
 1. SATA: limited by protocol, overhead of protocol, bandwidth of protocol
 2. PCIe: implement device driver, can do anything
 - a. example: FusionIO
 3. NVMe: standardized protocol over PCIe; very fast
- b. Flash Translation Layer
 - i. Implement SATA interface from disks
 1. read block, write block
 2. Manage internal parallelism
 - a. strip writes, reads across dies/planes
 - b. Wear leveling to handle endurance
 - ii. Internally implement on FLASH pages/blocks
 1. Challenge: doing writes
 - a. read-modify-write of whole erase block (cheap but slow and unreliable)
 - b. Log structuring (fast but complicated)
 2. With log structuring, need a map of where to find blocks
 - a. Compare to LFS inode map, but need for every block not just file inodes

- iii. Local map variations:
 - 1. Block-level only
 - a. can move a whole erase block, but pages must be at the right place in the erase block
 - b. Allows wear leveling, but still inefficient
 - c. Uses little memory
 - 2. Page-level
 - a. Complete flexibility (like virtual memory)
 - b. Need a lot of memory to store, or need extra reads to get from Flash (16x more space)
 - 3. Hybrid:
 - a. "Log area" for incoming writes (LFS active segments) that is page level
 - b. "Data area" for post-garbage collection data that is block level (plus bitmap of missing pages)
 - iv. Wear leveling:
 - 1. Dynamic wear leveling: change where data is written so not written to same place
 - 2. Static wear leveling: relocate data that hasn't been written to make low-wear spaces available
 - v. Garbage collection/cleaning
 - 1. Essentially same tradeoffs as LFS
 - 2. Most SSDs use greedy policy, not age-based policy
 - 3. Overprovision (like reserving space in LFS) to do less frequent GC
 - a. More blocks overwritten, so less data to be copied
 - b. Allows more GC in parallel, so GC takes less time
 - 4. Sustained write performance is the performance of the GC
- c. Persistence
 - i. Both data and logical map must be persistent to find data
 - ii. Cheap SSDs write logical map at shut down; on crash, may not have updated
 - 1. An extra metadata write per block write drops perf in half
 - iii. Solution: flash includes "out-of-band" area; 128 bytes per page
 - 1. Store reverse-map here
 - 2. Can reconstruct mapping by scanning these
 - 3. Reduce scanning by tracking which pages could be dirty

- d. SATA extensions
 - i. One new command needed: "Discard" or "Trim"
 - 1. Tells device when blocks are no longer needed by FS
 - 2. Need not be copied as part of GC
 - 3. Makes GC more like LFS – device is not "always full"
 - 4. Often slow – some SSDs erase data when discarded.
 - a. Much cheaper when larger
 - e. Performance take-aways
 - i. Random reads are fine, but a bit slower than sequential due to less parallelism
 - ii. Sequential writes tend to be faster, as don't need to GC if overwrite complete block
 - iii. Random writes devolve to speed of GC; if log-structure with page-based FTL can be as fast as sequential writes until GC kicks off
6. FlashTier – see slides